PCML Cheat Sheet

1 Math Prerequisites

– Bayes rule

$$p(A, B) = \underbrace{p(A|B)}_{\text{Lik.}} \underbrace{p(B)}_{\text{Prior}} = \underbrace{p(B|A)}_{\text{Post}} \underbrace{p(A)}_{\text{Marg. Lik.}}$$

– Gaussian distribution

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp(-\frac{(x-\mu)^2}{2\sigma^2})$$

$$\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{2\pi|\boldsymbol{\Sigma}|}} \exp(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}))$$

– Production of independent variables:

$$V(XY) = E(X^2)E(Y^2) - [E(X)]^2[E(Y)]^2$$

– Log-properties

$$\log(mn) = \log(m) + \log(n)$$
$$\log(m^n) = n\log(m)$$

– Covariance matrix of a data vector $\mathbf{x}_n$

$$\mathbf{C} = \frac{1}{N}\sum_{n=1}^{N}(\mathbf{x}_n - E[\mathbf{x}])(\mathbf{x}_n - E[\mathbf{x}])^T$$

1.1 Convexity

– A function is convex when a line joining two points never interects with the function anywhere else.
– A function $f(x)$ is convex, if for any $x_1, x_2 \in \mathbf{X}$ and for any $0 \le \lambda \le 1$, we have :

$$f(\lambda x_1 + (1-\lambda)x_2) \le \lambda f(x_1) + (1-\lambda)f(x_2)$$

– A function is strictly convex if the inequality is strict.
– A convex function has only one global minimum.
– Sums of convex functions are also convex.
– The Hessian is related to the convexity of a function: a twice differentiable function is convex if and only if the Hessian is positive definite.
– The Hessian of a convex function is positive semi-definite and for a strictly-convex function it is positive definite.
– The Hessian matrix of a function

$$\mathbf{H}_{i,j} = \frac{d^2 f}{dx_i dx_j}$$

1.2 Linear Algebra

– Column $\mathbf{x} \in R^n$, rows $\mathbf{x}^T$, matrix $\mathbf{A} \in R^{m \times n}$
– $\mathbf{x}^T\mathbf{x}$ is a scalar, $\mathbf{x}\mathbf{x}^T$ is a matrix
– $\mathbf{A}^{-1}$ exist if $\mathbf{A}$ is full rank
– **Condition number** of a function measures how much the output value can change for a small change in the input argument. A matrix with a high condition number is said to be **ill-conditioned**. If $\mathbf{A}$ is normal then

$$k(\mathbf{A}) = \left| \frac{\lambda_{max}(\mathbf{A})}{\lambda_{min}(\mathbf{A})} \right|$$

– A positive definite matrix is **symmetric** with all positive eigenvalues
– The real symmetric $N \times N$ matrix $\mathbf{V}$ is said to be **positive semidefinite** if

$$\mathbf{a}^T\mathbf{Va} \ge 0$$

for any real $N \times 1$ vector a.
– The real symmetric $N \times N$ matrix $\mathbf{V}$ is said to be **positive definite** if

$$\mathbf{a}^T\mathbf{Va} > 0$$

for any real $N \times 1$ vector a.
– Cost of matrix inversion: $O(n^3) \to O(n^{2.372})$
– Cost of determinant computation using LU decomposition: $O(n^3)$

2 Cost functions

– Cost functions are used to learn parameters that explain the data well.
– It is essential to make sure that a global minimum exist $\to$ lower bounded

Mean square error (MSE):

$$MSE(\boldsymbol{\beta}) = \sum_{n=1}^{N}(y_n - f(\mathbf{x}_n))^2$$

– MSE is **convex** thus it has only one global minumum value.
– MSE is not good when outliers are present.

Mean Absolute Error (MAE):

$$MAE = \sum_{n=1}^{N}|y_n - f(\mathbf{x}_n)|$$

Huber loss

$$Huber = \begin{cases} \frac{1}{2}z^2 & , |z| \le \delta \\ \delta|z| - \frac{1}{2}\delta^2 & , |z| > \delta \end{cases}$$

– Huber loss is convex, differentiable, and also robust to outliers
– hard to set δ.

Tukey's bisquare loss

$$L(z) = \begin{cases} z(\delta^2 - z^2)^2 & , |z| < \delta \\ 0 & , |z| \ge \delta \end{cases}$$

Tukey's loss is non-convex, non-differentiable, but robust to outliers.

Hinge loss

$$Hinge = [1 - y_n f(\mathbf{x}_n)]_+$$

Logistic loss

$$Logistic = \log(1 - \exp(y_n f(\mathbf{x}_n)))$$

3 Regression

– **Data** consists of N pairs $(y_n, \mathbf{x}_n)$

1. y_n the n'th output
2. $\mathbf{x}_n$ is a vector of D inputs

– **Prediction**: predict the ouput for a new input vector.

- **Interpretation**: understand the effect of inputs on output.
- **Outliers** are data that are far away from most of the other examples.

3.1 Linear Regression

- Model that assume linear relationship between inputs and the ouput.

$$y_n \equiv f(\mathbf{x}_n)$$
$$:= \beta_0 + \beta_1 x_{n1} + \dots$$
$$= \mathbf{x}_n^T \boldsymbol{\beta}$$

with $\boldsymbol{\beta}$ the parameters of the model.
- Variance grows only linearly with dimensionality

3.2 Gradient Descent

- Gradient descent uses only first-order information and takes steps in the direction of the gradient
- Given a cost function $\mathcal{L}(\boldsymbol{\beta})$ we wish to find $\boldsymbol{\beta}$ that minimizes the cost:

$$\min_{\boldsymbol{\beta}} \mathcal{L}(\boldsymbol{\beta})$$

3.2.1 Grid search

- Compute the cost over a grid of M points to find the minimum
- Exponential Complexity $O(NDM^D)$
- Hard to find a good range of values

3.3 Batch Gradient Descent

- Take steps in the opposite direction of the gradient

$$\boldsymbol{\beta}^{(k+1)} \leftarrow \boldsymbol{\beta}^{(k)} - \alpha \frac{d\mathcal{L}(\boldsymbol{\beta}^{(k)})}{d\boldsymbol{\beta}}$$

with $\alpha > 0$ the learning rate.
- With α too big, method might diverge. With α too small, convergence is slow.

3.4 Gradients for MSE

$$\tilde{\mathbf{X}} = [1 \ \ \mathbf{X}]$$

- We define the error vector $\mathbf{e}$:

$$\mathbf{e} = \mathbf{y} - \tilde{\mathbf{X}}\boldsymbol{\beta}$$

- and MSE as follows:

$$\mathcal{L}(\boldsymbol{\beta}) = \frac{1}{2N}\sum_{n=1}^{N}(y_n - \tilde{\mathbf{x}}_n^T\boldsymbol{\beta})^2 = \frac{1}{2N}\mathbf{e}^T\mathbf{e}$$

- then the gradient is given by

$$\frac{d\mathcal{L}}{d\boldsymbol{\beta}} = -\frac{1}{N}\tilde{\mathbf{X}}^T\mathbf{e}$$

- Optimality conditions:

  1. *necessary*: gradient equal to zero: $\frac{d\mathcal{L}(\boldsymbol{\beta}^*)}{d\boldsymbol{\beta}} = 0$

  2. *sufficient*: Hessian matrix is positive definite: $\mathbf{H}(\boldsymbol{\beta}^*) = \frac{d^2\mathcal{L}(\boldsymbol{\beta}^*)}{d\boldsymbol{\beta}d\boldsymbol{\beta}^T}$

- Very sensitive to illconditioning. Therefore, always normalize your feature otherwise step-size selection is difficult since different directions might move at different speed.
- *Complexity*: $O(NDI)$ with I the number of iterations

3.5 Least Squares

- In some cases, we can compute the minimum of the cost function analytically.
- use the first optimality conditions:

$$\frac{d\mathcal{L}}{d\boldsymbol{\beta}} = 0 \Rightarrow \tilde{\mathbf{X}}^T\mathbf{e} = \tilde{\mathbf{X}}^T(\mathbf{y} - \tilde{\mathbf{X}}\boldsymbol{\beta}) = 0$$

- When $\tilde{\mathbf{X}}^T\tilde{\mathbf{X}}$ is invertible, we have the closed-form expression

$$\boldsymbol{\beta}^* = (\tilde{\mathbf{X}}^T\tilde{\mathbf{X}})^{-1}\tilde{\mathbf{X}}^T\mathbf{y}$$

- thus we can predict values for a new $\mathbf{x}_*$

$$y_* = \tilde{\mathbf{x}}_*^{\mathbf{T}}\boldsymbol{\beta}^* = \tilde{\mathbf{x}}_*^{\mathbf{T}}(\tilde{\mathbf{X}}^T\tilde{\mathbf{X}})^{-1}\tilde{\mathbf{X}}^T\mathbf{y}$$

- The **Gram matrix** $\tilde{\mathbf{X}}^T\tilde{\mathbf{X}}$ is positive definite and is also invertible iff $\tilde{\mathbf{X}}$ has full column rank.
- *Complexity*: $O(ND^2 + D^3) \equiv O(ND^2)$
- $\tilde{\mathbf{X}}$ can be rank deficient when $D > N$ or when the comlumns $\bar{\mathbf{x}}_d$ are nearly collinear. In this case, the matrix is ill-conditioned, leading to numerical issues.

3.6 Maximum Likelihood

- Let define our mistakes $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$.

$$\rightarrow y_n = \tilde{\mathbf{x}}_n^T\boldsymbol{\beta} + \epsilon_n$$

- Another way of expressing this:

$$p(\mathbf{y}|\tilde{\mathbf{X}}, \boldsymbol{\beta}) = \prod_{n=1}^{N} p(y_n|\tilde{\mathbf{x}}_n, \boldsymbol{\beta})$$
$$= \prod_{n=1}^{N} \mathcal{N}(y_n|\tilde{\mathbf{x}}_n^T\boldsymbol{\beta}, \sigma^2)$$

which defines the likelihood of observating $\mathbf{y}$ given $\tilde{\mathbf{X}}$ and $\boldsymbol{\beta}$
- Define cost with log-likelihood

$$\mathcal{L}_{lik}(\boldsymbol{\beta}) = \log p(\mathbf{y}|\tilde{\mathbf{X}}, \boldsymbol{\beta})$$
$$= -\frac{1}{2\sigma^2}\sum_{n=1}^{N}(y_n - \tilde{\mathbf{x}}_n^T\boldsymbol{\beta})^2 + cnst$$

- Maximum likelihood estimator (MLE) gives another way to design cost functions

$$\operatorname*{argmin}_{\boldsymbol{\beta}} \mathcal{L}_{mse}(\boldsymbol{\beta}) = \operatorname*{argmax}_{\boldsymbol{\beta}} \mathcal{L}_{lik}(\boldsymbol{\beta})$$

- MLE can also be interpreted as finding the model under which the obeserved data is most likely to have been generated from.
- With Laplace distribution

$$p(y_n|\tilde{\mathbf{x}}_n, \boldsymbol{\beta}) = \frac{1}{2b}e^{-\frac{1}{b}|y_n - \tilde{\mathbf{x}}_n^T\boldsymbol{\beta}|}$$

$$\sum_n \log p(y_n|\tilde{\mathbf{x}}_n, \boldsymbol{\beta}) = \sum_n |y_n - \tilde{\mathbf{x}}_n^T\boldsymbol{\beta}| + cnst$$

3.7 Ridge Regression

– Linear models usually underfit. One way is to use nonlinear basis functions instead.

$$y_n = \beta_0 + \sum_{j=1}^{M} \beta_j \phi_j(\mathbf{x}_n) = \tilde{\phi}(\mathbf{x}_n)^T \boldsymbol{\beta}$$

– This model is linear in $\boldsymbol{\beta}$ but nonlinear in $\mathbf{x}$. Note that the dimensionality is now M, not D.
– Polynomial basis

$$\phi(x_n) = [1, x_n, x_n^2, ..., x_n^M]$$

– The least square solution becomes

$$\boldsymbol{\beta}_{lse}^* = (\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}})^{-1} \tilde{\boldsymbol{\Phi}}^T \mathbf{y}$$

– Complex models overfit easily. Thus we can choose simpler models by adding a **regularization term** which penalizes complex models

$$\min_{\boldsymbol{\beta}} \left(\mathcal{L}(\boldsymbol{\beta}) + \frac{\lambda}{2N} \sum_{j=1}^{M} \beta_j^2 \right)$$

$$\boldsymbol{\beta}^* = \underset{\boldsymbol{\beta}}{\operatorname{argmin}} \left(\frac{1}{2} (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) + \frac{\lambda}{2} \boldsymbol{\beta}^T \boldsymbol{\beta} \right)$$

– Note that β_0 is not penalized.
– By differentiating and setting to zero we get

$$\boldsymbol{\beta}_{ridge} = (\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}} + \boldsymbol{\Lambda})^{-1} \tilde{\boldsymbol{\Phi}}^T \mathbf{y}$$

$$\boldsymbol{\Lambda} = \begin{bmatrix} 0 & \underline{0} \\ \underline{0} & \lambda I_m \end{bmatrix}$$

– Ridge regression improves the condition number of the Gram matrix since the eigenvalues of $(\tilde{\boldsymbol{\Phi}}^T \tilde{\boldsymbol{\Phi}} + \lambda I_m)$ are at least λ
– **Maximum-a-posteriori (MAP) estimator**:

   – Maximizes the product of the likelihood *and* the **prior**.

$$\boldsymbol{\beta}_{map} = \underset{\boldsymbol{\beta}}{\operatorname{argmax}} \left(p(\mathbf{y}|\mathbf{X}, \boldsymbol{\Lambda}) p(\boldsymbol{\beta}|\boldsymbol{\Sigma}) \right)$$

   – Assume $\beta_0 = 0$

$$\boldsymbol{\beta}_{ridge} = \underset{\boldsymbol{\beta}}{\operatorname{argmax}} \left(\log \left[\prod_{n=1}^{N} \mathcal{N}(y_n|\mathbf{x}_n^T\boldsymbol{\beta}, \boldsymbol{\Lambda}) \times \mathcal{N}(\boldsymbol{\beta}|0, \mathbf{I}) \right] \right)$$

– **Lasso regularizer** forces some β_i to be strictly 0 and therefore forces sparsity in the model.

$$\min_{\boldsymbol{\beta}} \frac{1}{2N} \sum_{n=1}^{N} (y_n - \tilde{\phi}(\mathbf{x}_n)^T \boldsymbol{\beta})^2, \quad \text{such that} \sum_{i=1}^{M} |\beta_i| \leq \tau$$

3.8 Cross-Validation

– We should choose λ to minimize the mistakes that will be made in the future.

– We split the data into train and validation sets and we pretend that the validation set is the future data. We fit our model on the training set and compute a prediction-error on the validation set. This gives us an *estimate* of the *generalization error*.
– **K-fold cross validation** randomly partition the data into K groups. We train on $K-1$ groups and test on the remaining group. We repeat this until we have tested on all K sets. We then average the results.
– Cross-validation returns an unbiased estimate of the generalization error and its variance.

3.9 Bias-Variance decomposition

– The expected test error can be expressed as the sum of two terms

   – **Squared bias**: The average *shift* of the predictions
   – **Variance**: measure how data points vary around their average.

$$\text{expected loss} = (\text{bias})^2 + \text{variance} + \text{noise}$$

– Both model bias and estimation bias are important
– Ridge regression increases estimation bias while reducing variance
– Increasing model complexity increases test error

$$\text{Small } \lambda \rightarrow \text{low bias but large variance}$$

$$\text{Large } \lambda \rightarrow \text{large bias but low variance}$$

3.10 Logistic Regression

– **Classification** relates input variables $\mathbf{x}$ to discrete output variable y
– **Binary classifier**: we use $y = 0$ for $\mathbf{C}_1$ and $y = 1$ for $\mathbf{C}_2$.
– Can use least-squares to predict $\hat{y}_*$

$$\hat{y} = \begin{cases} \mathbf{C}_1 & \hat{y}_* < 0.5 \\ \mathbf{C}_2 & \hat{y}_* \geq 0.5 \end{cases}$$

– **Logistic function**

$$\sigma(x) = \frac{\exp(x)}{1 + \exp(x)}$$

$$p(y_n = \mathbf{C}_1|\mathbf{x}_n) = \sigma(\tilde{\mathbf{x}}^T \boldsymbol{\beta})$$
$$p(y_n = \mathbf{C}_2|\mathbf{x}_n) = 1 - \sigma(\tilde{\mathbf{x}}^T \boldsymbol{\beta})$$

– The probabilistic model:

$$p(\mathbf{y}|\mathbf{X}, \boldsymbol{\beta}) = \prod_{n=1}^{N} \sigma(\tilde{\mathbf{x}}_n^T \boldsymbol{\beta})^{y_n} (1 - \sigma(\tilde{\mathbf{x}}_n^T \boldsymbol{\beta})^{1-y_n}$$

– The log-likelihood:

$$\mathcal{L}_{mle}(\boldsymbol{\beta}) = \sum_{n=1}^{N} \left(y_n \tilde{\mathbf{x}}_n^T \boldsymbol{\beta} - \log(1 + \exp(\tilde{\mathbf{x}}_n^T \boldsymbol{\beta})) \right)$$

– We can use the fact that

$$\frac{d}{dx} \log(1 + \exp(x)) = \sigma(x)$$

– Gradient of the log-likelihood

$$\mathbf{g} = \frac{d\mathcal{L}}{d\boldsymbol{\beta}} = \sum_{n=1}^{N} \left(\tilde{\mathbf{x}}_n y_n - \tilde{\mathbf{x}}_n \sigma(\tilde{\mathbf{x}}_n^T \boldsymbol{\beta}) \right)$$

$$= -\tilde{\mathbf{X}}^T [\sigma(\tilde{\mathbf{X}} \boldsymbol{\beta}) - \mathbf{y}]$$

– The negative of the log-likelihood $-\mathcal{L}_{mle}(\boldsymbol{\beta})$ is convex
– **Hessian** of the log-likelihood
  – We know that

$$\frac{d\sigma(t)}{dt} = \sigma(t)(1 - \sigma(t))$$

  – Hessian is the derivative of the gradient

$$\mathbf{H}(\boldsymbol{\beta}) = -\frac{d\mathbf{g}(\boldsymbol{\beta})}{d\boldsymbol{\beta}^T} = \sum_{n=1}^{N} \frac{d}{d\boldsymbol{\beta}^T} \sigma(\tilde{\mathbf{x}}_n^T \boldsymbol{\beta}) \tilde{\mathbf{x}}_n$$

$$= \sum_{n=1}^{N} \tilde{\mathbf{x}}_n \sigma(\tilde{\mathbf{x}}_n^T \boldsymbol{\beta})(1 - \sigma(\tilde{\mathbf{x}}_n^T \boldsymbol{\beta})) \tilde{\mathbf{x}}_n^T$$

$$= \tilde{\mathbf{X}}^T \mathbf{S} \tilde{\mathbf{X}}$$

  where $\mathbf{S}$ is a $N \times N$ diagonal matrix with diagonals

$$S_{nn} = \sigma(\tilde{\mathbf{x}}_n^T \boldsymbol{\beta})(1 - \sigma(\tilde{\mathbf{x}}_n^T \boldsymbol{\beta}))$$

  – The negative of the log-likelihood is not strictly convex.
– **Newton's Method**
  – Uses second-order information and takes steps in the direction that minimizes a quadratic approximation

$$\mathcal{L}(\boldsymbol{\beta}) = \mathcal{L}(\boldsymbol{\beta}^{(k)}) + \mathbf{g}_k^T(\boldsymbol{\beta} - \boldsymbol{\beta}^{(k)}) + (\boldsymbol{\beta} - \boldsymbol{\beta}^{(k)})^T \mathbf{H}_k (\boldsymbol{\beta} - \boldsymbol{\beta}^{(k)})$$

  and it's minimum is at

$$\boldsymbol{\beta}^{k+1} = \boldsymbol{\beta}^{(k)} - \alpha_k \mathbf{H}_k^{-1} \mathbf{g}_k$$

  where $\mathbf{g}_k$ is the gradient and α_k the learning rate.
  – Complexity: $O((ND^2 + D^3)I)$
– **Penalized Logistic Regression**

$$\min_{\boldsymbol{\beta}} \left(-\sum_{n=1}^{N} \log p(y_n | \mathbf{x}_n^T \boldsymbol{\beta}) + \lambda \sum_{d=1}^{D} \beta_d^2 \right)$$

4 Generalized Linear Model

– **Exponential family distribution**

$$p(\mathbf{y}|\boldsymbol{\eta}) = \frac{h(y)}{Z} \exp(\boldsymbol{\eta}^T \boldsymbol{\phi}(\mathbf{y}) - A(\boldsymbol{\eta}))$$

– Bernoulli distribution

$$p(y|\mu) = \mu^y (1 - \mu)^{1-y}$$

$$= \exp(y \log(\frac{\mu}{1 - \mu} + \log(1 - \mu)))$$

– there is a relationship between η and μ throught the **link function**

$$\eta = \log(\frac{\mu}{1 - \mu}) \leftrightarrow \mu = \frac{e^\eta}{1 + e^\eta}$$

– Note that μ is the mean parameter of y
– Relationship between the mean $\boldsymbol{\mu}$ and $\boldsymbol{\eta}$ is defined using a link function g

$$\boldsymbol{\eta} = \mathbf{g}(\boldsymbol{\mu}) \Leftrightarrow \boldsymbol{\mu} = \mathbf{g}^{-1}(\boldsymbol{\eta})$$

– First and second derivatives of $A(\eta)$ are related to the mean and the variance

$$\frac{dA(\eta)}{d\eta} = E[\phi(\eta)], \quad \frac{d^2 A(\eta)}{d\eta^2} = Var[\phi(\eta)]$$

– $A(\eta)$ is convex
– The generalized maximum likelihood cost to minimize is

$$\min_{\boldsymbol{\beta}} \mathcal{L}(\boldsymbol{\beta}) = -\sum_{n=1}^{N} \log(p(y_n | \tilde{\mathbf{x}}_n^T \boldsymbol{\beta}))$$

where $p(y_n | \tilde{\mathbf{x}}_n^T \boldsymbol{\beta})$ is an exponential family distribution
– We obtain the solution

$$\frac{d\mathcal{L}}{d\boldsymbol{\beta}} = \tilde{\mathbf{X}}^T [\mathbf{g}^{-1}(\boldsymbol{\eta}) - \boldsymbol{\phi}(\mathbf{y})]$$

5 k-Nearest Neighbor (k-NN)

– The k-NN prediction for $\mathbf{x}$ is

$$f(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_n \in nbh_k(\mathbf{x})} y_n$$

where $nbh_k(\mathbf{x})$ is the neightborhood of $\mathbf{x}$ defined by the k closest points $\mathbf{x}_n$ in the training data
– **Curse of dimensionality**: Generalizing correctly becomes exponentially harder as the dimensionality grows.
– Gathering more inputs variables may be a bad thing

6 Kernel Ridge Regression

– The following is true for ridge regression

$$\boldsymbol{\beta} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}_D)^{-1} \mathbf{X}^T \mathbf{y} \qquad (1)$$

$$= \mathbf{X}^T (\mathbf{X}\mathbf{X}^T + \lambda \mathbf{I}_N)^{-1} \mathbf{y} \qquad (2)$$

$$= \mathbf{X}^T \boldsymbol{\alpha}$$

– Complexity of computing $\boldsymbol{\beta}$

1. $O(D^2 N + D^3)$
2. $O(N^2 D + N^3)$

– Thus we have

$$\boldsymbol{\beta} = \sum_{n=1}^{N} \alpha_n \mathbf{x}_n, \quad \mathbf{y} = \sum_{d=1}^{D} \beta_d \bar{\mathbf{x}}_d$$

with $\mathbf{x}_n$ the rows of $\mathbf{X}$ and $\bar{\mathbf{x}}_d$ the columns of $\mathbf{X}$
– The representer theorem allows us to write an equivalent optimization problem n terms of $\boldsymbol{\alpha}$.

$$\boldsymbol{\alpha} = \underset{\boldsymbol{\alpha}}{\mathrm{argmax}} \left(-\frac{1}{2}\boldsymbol{\alpha}(\mathbf{X}\mathbf{X}^T + \lambda \mathbf{I}_N)^T \boldsymbol{\alpha} + \boldsymbol{\alpha}^T \mathbf{y} \right)$$

– $\mathbf{X}\mathbf{X}^T$ is called the **kernel matrix**
– Kernelized rigde regression might be computationally more efficient in some cases.
– **Kernel trick**:
  – We can work directly with $\mathbf{K} = \mathbf{X}\mathbf{X}^T$ and never have to worry about $\mathbf{X}$
  – Using the basis function $\boldsymbol{\phi}(\mathbf{x})$, we do not need to specify it explicitly, since we can work directly with $\mathbf{K} = \boldsymbol{\phi}(\mathbf{x})$
  – We will use a kernel function $k(\mathbf{x}, \mathbf{x}') = \boldsymbol{\phi}(\mathbf{x})^T \boldsymbol{\phi}(\mathbf{x}')$
  – The evaluation of a kernel is usually faster with k than with $\boldsymbol{\phi}$
– **Radial Basis function kernel (RBF)**

$$k(\mathbf{x}, \mathbf{x}') = \exp(-\frac{1}{2}(\mathbf{x} - \mathbf{x}')^T (\mathbf{x} - \mathbf{x}'))$$

– Properties of a kernel to ensure the existance of a corresponding $\boldsymbol{\phi}$:
  – $\mathbf{K}$ should be symmetric: $k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$
  – $\mathbf{K}$ should be positive semidefinite.

7 Support Vector Machine

– **Hinge loss**

$$[t]_+ = \max(0, t)$$

– Solution to the dual problem is sparse and non-zero entries will be our **support vectors**.
– Assume $y_n \in \{-1, 1\}$
– SVM optimizes the following cost

$$\min_{\boldsymbol{\beta}} \sum_{n=1}^{N} [1 - y_n \tilde{\boldsymbol{\phi}}_n^T \boldsymbol{\beta}]_+ + \frac{\lambda}{2} \sum_{j=1}^{M} \beta_j^2$$

– The minimum doesn't change with a rescaling of $\boldsymbol{\beta}$
– **Duality**:
  – Hard to minimize $g(\boldsymbol{\beta})$ so we define

$$g(\boldsymbol{\beta}) = \max_{\boldsymbol{\alpha}} G(\boldsymbol{\beta}, \boldsymbol{\alpha})$$

  – we use the property that

$$C[v_n]_+ = \max(0, Cv_n) = \max_{\alpha_n \in [0,C]} \alpha_n v_n$$

  – We can rewrite the problem as

$$\min_{\boldsymbol{\beta}} \max_{\boldsymbol{\alpha} \in [0,C]^N} \sum_{n=1}^{N} \alpha_n (1 - y_n \tilde{\boldsymbol{\phi}}_n^T \boldsymbol{\beta}) + \frac{1}{2} \sum_{j=1}^{M} \beta_j^2$$

– This is differentiable, convex in $\boldsymbol{\beta}$ and concave in $\boldsymbol{\alpha}$
– **Minimax theorem**:

$$\min_{\boldsymbol{\beta}} \max_{\boldsymbol{\alpha}} G(\boldsymbol{\beta}, \boldsymbol{\alpha}) = \max_{\boldsymbol{\alpha}} \min_{\boldsymbol{\beta}} G(\boldsymbol{\beta}, \boldsymbol{\alpha})$$

because G is convex in $\boldsymbol{\beta}$ and concave in $\boldsymbol{\alpha}$.
– Derivative w.r.t. $\boldsymbol{\beta}$:

$$\frac{dG}{d\boldsymbol{\beta}} = -\left(\sum_{n=1}^{N} \alpha y_n \tilde{\phi}_n \right) + \begin{bmatrix} 0 \\ \boldsymbol{\beta}_{1:M} \end{bmatrix}$$

– Equating this to 0, we get:

$$\boldsymbol{\beta}_{1:M}^* = \sum_{n=1}^{N} \alpha_n y_n \phi_n = \boldsymbol{\Phi}^T \mathrm{diag}(\mathbf{y}) \boldsymbol{\alpha}$$

$$\boldsymbol{\alpha}^T \mathbf{y} = 0$$

– Plugging $\boldsymbol{\beta}^*$ back in the dual problem

$$\max_{\boldsymbol{\alpha} \in [0,C]^N} \boldsymbol{\alpha}^T \mathbf{1} - \frac{1}{2} \boldsymbol{\alpha}^T \mathbf{Y} \boldsymbol{\Phi} \boldsymbol{\Phi} \mathbf{Y} \boldsymbol{\alpha}$$

– This is a differentiable least-squares problem. Optimization is easy using Sequential Minimal Optimization. It is also naturally kernelized with $\mathbf{K} = \boldsymbol{\Phi}\boldsymbol{\Phi}^T$
– The solution $\boldsymbol{\alpha}$ is sparse and is non-zero only for the training examples that are instrumental in determining the decision boundary.

8 K-means

– **Unsupervised learning**: Represent particular input patterns in a way that reflects the statistical structure of the overall collections of input partterns.
– **Cluster** are groups of points whose inter-point distances are small compared to the distances outside the cluster.

$$\min_{\mathbf{r}, \boldsymbol{\mu}} \mathcal{L}(\mathbf{r}, \boldsymbol{\mu}) = \sum_{k=1}^{K} \sum_{n=1}^{N} r_{nk} ||\mathbf{x}_n - \boldsymbol{\mu}_k||_2^2$$

such that $r_{nk} \in \{0, 1\}$ and $\sum_{k=1}^{K} r_{nk} = 1$
– K-means algorithm:
Initialize $\boldsymbol{\mu}_k$, then iterate

1. For all n, compute $\mathbf{r}_n$ given $\boldsymbol{\mu}$

$$r_{nk} = \begin{cases} 1 & \text{if } k = \mathrm{argmin}_j ||\mathbf{x}_n - \boldsymbol{\mu}||_2^2 \\ 0 & \text{otherwise} \end{cases}$$

2. For all k, compute μ_k given $\mathbf{r}$

$$\boldsymbol{\mu}_k = \frac{\sum_{n=1}^{N} r_{nk} \mathbf{x}_n}{\sum_{n=1}^{N} r_{nk}}$$

– A good initialization procedure is to choose the prototypes to be equal to a random subset of K data points.

– Probabilistic model

$$p(\mathbf{r}, \boldsymbol{\mu}) = \prod_{n=1}^{N} \prod_{k=1}^{K} \left[\mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \mathbf{I}) \right]^{r_{nk}}$$

– Computation can be heavy, each example can belong to only on cluster and clusters have to be spherical.

9 Gaussian Mixture Models

– Clusters can be spherical using a full covariance matrix instead of isotropic covariance.

$$p(\mathbf{X} | \boldsymbol{\mu}, \boldsymbol{\Sigma}, \mathbf{r}) = \prod_{n=1}^{N} \prod_{k=1}^{K} \left[\mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \Sigma_k) \right]^{r_{nk}}$$

– **Soft-clustering**: Points can belong to several cluster by defining r_n to be a random variable.

$$p(r_{nk} = 1) = \pi_k \text{ where } \pi_k > 0, \forall k$$

$$\sum_{k=1}^{K} \pi_k = 1$$

– Joint distribution of Gaussian mixture model

$$p(\mathbf{X}, \mathbf{r} | \boldsymbol{\mu}, \boldsymbol{\Sigma}, \boldsymbol{\pi}) = \prod_{n=1}^{N} \left[p(\mathbf{x}_n | r_n, \boldsymbol{\mu}, \Sigma) p(r_n | \boldsymbol{\pi}) \right]$$

$$= \left[\prod_{k=1}^{K} [(\mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \Sigma_k))^{r_{nk}} \prod_{k=1}^{K} [\pi]^{r_{nk}} \right]$$

$$\text{Joint} = \text{Likelihood} \times \text{Prior}$$

– r_n are called *latent* unobserved variables
– Unknown parameters are given by $\boldsymbol{\theta} = \{\boldsymbol{\mu}, \boldsymbol{\Sigma}, \boldsymbol{\pi}\}$
– We get the **marginal likelihood** by marginalizing r_n out from the likelihood

$$p(\mathbf{x}_n | \boldsymbol{\theta}) = \sum_{k=1}^{K} p(\mathbf{x}_n, r_n = k | \boldsymbol{\theta})$$

$$= \sum_{k=1}^{K} p(r_n = k | \boldsymbol{\theta}) p(\mathbf{x}_n | r_n = k, \boldsymbol{\theta})$$

$$= \sum_{k=1}^{K} \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

– Without a latent variable model, number of parameters grow at rate $O(N)$
– After marginalization, the growth is reduced to $O(D^2 K)$
– To get maximum likelihood estimate of $\boldsymbol{\theta}$, we maximize

$$\max_{\boldsymbol{\theta}} \sum_{n=1}^{N} \log \sum_{k=1}^{K} \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\sigma}_k)$$

10 Expectation Maximization Algorithm

– *[ALGORITHM]* Start with $\boldsymbol{\theta}^{(1)}$ and iterate

1. *Expectation step*: Compute a lower bound to the cost such that it is tight at the previous $\boldsymbol{\theta}^{(i)}$

$$\log \sum_{k=1}^{K} \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \geq \sum_{k=1}^{K} p_{kn} \log \frac{\pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{p_{kn}}$$

with equality when,

$$p_{kn} = \frac{\pi_k \mathcal{N}(\mathbf{x}_n, \boldsymbol{\Sigma}_k)}{\sum_{k=1}^{K} \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}$$

2. *Maximization step*: Update $\boldsymbol{\theta}$

$$\boldsymbol{\theta}^{(i+1)} = \underset{\boldsymbol{\theta}}{\text{argmax}} \, \mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\theta}^{(i)})$$

$$\boldsymbol{\mu}_k^{i+1} = \frac{\sum_{n=1}^{N} p_{kn}^{(i)} \mathbf{x}_n}{\sum_{n=1}^{N} p_{kn}^{(i)}}$$

$$\boldsymbol{\Sigma}_k^{(i+1)} = \frac{\sum_{n=1}^{N} p_{kn}^{(i)} (\mathbf{x}_n - \boldsymbol{\mu}_k^{(i+1)})(\mathbf{x}_n - \boldsymbol{\mu}_k^{(i+1)})^T}{\sum_{n=1}^{N} p_{kn}^{(i)}}$$

$$\pi_k^{(i+1)} = \frac{1}{N} \sum_{n=1}^{N} p_{kn}^{(i)}$$

– If the covariance is diagonal, then we have K-means.

11 Matrix factorization

– We have D movies and N users
– $\mathbf{X}$ is a matrix $D \times N$ with x_{dn} the rating of n'th user for d'th movie.
– We project data vectors $\mathbf{x}_n$ to a smaller dimension $\mathbf{z}_n \in \mathbb{R}^M$
– We have now 2 latent variables:

  – $\mathbf{Z}$ a $N \times M$ matrix that gives features for the users
  – $\mathbf{W}$ a $D \times M$ matrix that gives features for the movies

$$x_{dn} \approx \mathbf{w}_d^T \mathbf{z}_n$$

– We can add a regularizer and minimize the following cost:

$$\mathcal{L}(\mathbf{W}, \mathbf{Z}) = \frac{1}{2} \sum_{n=1}^{N} \sum_{d=1}^{D} (x_{dn} - \mathbf{w}_d^T \mathbf{z}_n)^2 + \frac{\lambda_w}{2} \sum_{d=1}^{D} \mathbf{w}_d^T \mathbf{w}_d + \frac{\lambda_z}{2} \sum_{n=1}^{N} \mathbf{z}_n^T \mathbf{z}_n$$

– We can use coordinate descent algorithm, by first minimizing w.r.t. $\mathbf{Z}$ given $\mathbf{W}$ and then minimizing $\mathbf{W}$ given $\mathbf{Z}$. This is called **Alternating least-squares (ALS)**:

$$\mathbf{Z}^T \leftarrow (\mathbf{W}^T \mathbf{W} + \lambda_z \mathbf{I}_M)^{-1} \mathbf{W}^T \mathbf{X}$$
$$\mathbf{W}^T \leftarrow (\mathbf{Z}^T \mathbf{Z} + \lambda_w \mathbf{I}_M)^{-1} \mathbf{Z}^T \mathbf{X}^T$$

– *Complexity*: $O(DNM^2 + NM^3) \rightarrow O(DNM^2)$

– Probabilistic model

$$\prod_{n=1}^{N} \prod_{d \in O_n} \mathcal{N}(x_{dn} | \mathbf{w}_d^T \mathbf{z}_n, I) \times \prod_{n=1}^{N} \mathcal{N}(\mathbf{z}_n | 0, \frac{1}{\lambda_z} I) \times \prod_{d=1}^{D} \mathcal{N}(\mathbf{w}_d | 0, \frac{1}{\lambda_w} I)$$

– Since many ratings are missing we cannot normalize the data. A solution is to add offset terms:

$$\frac{1}{2} \sum_{n=1}^{N} \sum_{d \in O_n} (x_{dn} - \mathbf{w}_d^T \mathbf{z}_n - w_{0d} - z_{0n} - \mu)^2$$

12 Singular Value Decomposition

– Matrix factorization method

$$\mathbf{X} = \mathbf{U}\mathbf{S}\mathbf{V}^T$$

– $\mathbf{U}$ is an $D \times D$ matrix
– $\mathbf{V}$ is an $N \times N$ matrix
– $\mathbf{S}$ is a non-negative diagonal matrix of size $D \times N$ which are called **singular values** appearing in a descending order.
– Columns of $\mathbf{U}$ and $\mathbf{V}$ are the left and right **singular vectors** respectively.
– Assuming $D < N$ we have

$$\mathbf{X} = \sum_{d=i}^{D} s_d \mathbf{u}_d \mathbf{v}_d^T$$

This tells you about the spectrum of $\mathbf{X}$ where higher singular vectors contain the *low-frequency information* and lower singular values contain the *high-frequency information*.
– Let's now truncate these matrices

$$\mathbf{X} \approx \mathbf{U}_{tr} \mathbf{S}_{tr} \mathbf{V}_{tr}^*$$
$$\mathbf{T}_{tr} \approx \mathbf{U}_{tr} \mathbf{S}_{tr}$$
$$\mathbf{T}_{te} \approx \mathbf{X}_{te} \mathbf{V}_{tr}$$

with $\mathbf{T}_{tr}$ the reduced feature set of $\mathbf{X}$ and $\mathbf{T}_{te}$ the reduced feature set of $\mathbf{X}_{te}$

13 Principal Componence Analysis

– PCA is a dimensionality reduction method and a method to decorrelate the data

$$\mathbf{X} \approx \tilde{\mathbf{X}} = \mathbf{W}\mathbf{Z}^T$$

such that columns of $\mathbf{W}$ are orthogonal.
– If the data is zero mean

$$\mathbf{C} = \frac{1}{N} \mathbf{X}\mathbf{X}^T \Rightarrow \mathbf{X}\mathbf{X}^T = \mathbf{U}\mathbf{S}^2\mathbf{U}^T$$
$$\Rightarrow \mathbf{U}^T \mathbf{X}\mathbf{X}^T \mathbf{U} = \mathbf{U}^T \mathbf{U}\mathbf{S}^2\mathbf{U}^T \mathbf{U} = \mathbf{S}^2$$

– Thus the columns of matrix $\mathbf{U}$ are called the **principal components** and they decorrelate the covariance matrix.
– Using SVD, we can compute the matrices in the following way

$$\mathbf{W} = \mathbf{U}\mathbf{S}^{1/2}$$
$$\mathbf{Z} = \mathbf{V}\mathbf{S}^{1/2}$$

14 Belief Propagation

– the goal is to learn *inference of the latent variables using belief propagation*
– Given a directed acyclic graph G and parameters θ, a **Bayesian network** defines the joint distribution as follows

$$p_\theta(\mathbf{x}) = \prod_{k=1}^{K} p_\theta(x_k | \text{parents}_k)$$

– We can reduce the complexity by using the structure of the problem and the bayesian rule.

15 Multi-Layer Perceptron (MLP)

– Known as **feed-forward neural network**
– $\mathbf{z}_n^{(k)}$ is the k'th hidden vector
– $\mathbf{a}_n^{(k)}$ is the corresponding activation
– There are a total of K layers

$$a_{mn}^{(k)} = (\boldsymbol{\beta}_m^{(k)})^T \mathbf{z}_n^{(k-1)}, \quad z_{mn}^{(k)} = h(a_{mn}^{(k)})$$

– For the first layer we have $\mathbf{z}_n^{(0)} = \mathbf{x}_n$
– For the last layer, we use a link function to map $\mathbf{z}_n^{(K-1)}$ to the output $\mathbf{y}_n$
– A 1-layer MLP is simply a generalization of linear/logistic regression
– $\mathbf{B}^{(k)}$ a matrix with rows $(\boldsymbol{\beta}_m^{(k)})^T$

$$\mathbf{a}_n^{(k)} = \mathbf{B}^{(k)} \mathbf{z}_n^{(k-1)}, \quad \mathbf{z}_n^{(k)} = h(\mathbf{a}_n^{(k)})$$

– thus we have the input-output relationship

$$\hat{y}_n = g((\boldsymbol{\beta}^{(K-1)})^T * h(\mathbf{B}^{(K-2)} * h(* ... * h(\mathbf{B}^{(1)} * \mathbf{x}_n))))$$

with $g()$ the link function
– We learn parameters $\mathbf{B}$ using stochastic gradient-descent
– Frequently used transfer function

$$g(a) = \tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$

– **Backpropagation** is a technic to compute the gradient in time linear in the number of training points and the number of weights.

16 Gaussian Process (GP)

– GP are a family of statistical distributions in which time plays a role and for which any finite linear combination of samples has a joint Gaussian distribution.
– They can be completely defined by their second-order statistics which means that if they have mean zero, then defining the covariance function completely defines the process behaviour.
– Let us place a probabilistic prior shape on the approximation of a function.
– A GP process defines a prior over function f

$$p(\mathbf{f}|X) = \mathcal{N}(\mathbf{f}|\mathbf{0}, K(X))$$

– $K(X)$ defines shape and prior knowledge about our problem

$$p(\mathbf{f}|\mathbf{y}, \mathbf{X}_*, \mathbf{X}) \sim \mathcal{N}(\mu', \sigma')$$
$$\mu' = K(\mathbf{X}_*, \mathbf{X})[K(\mathbf{X}, \mathbf{X}) + \sigma_n^2 I]^{-1} \mathbf{y}$$
$$\sigma' = K(\mathbf{X}_*, \mathbf{X}_*) - K(\mathbf{X}_*, \mathbf{X})[K(\mathbf{X}, \mathbf{X}) + \sigma_n^2 I] K(\mathbf{X}, \mathbf{X}_*)$$

with σ_n the variance of the noise
– RBF kernel

$$k(\mathbf{x}_n, \mathbf{x}_m) = e^{-||\mathbf{x}_n - \mathbf{x}_m||^2 / L^2}$$

– Quadratic kernel

$$k(\mathbf{x}_n, \mathbf{x}_m) = (1 + \mathbf{x}_n^T \mathbf{x}_m)^2$$

17 Decision Trees (DT)

– A decision tree is a structure in which each internal node represents a test on an attribute and each branch represents the outcome of the test and each leaf represents a class label.
– Fast to train and fast to make predictions
– Efficient for very high dimensional feature spaces and very large amounts of training data
– Lack of smoothness and high variance (overfitting)
– Goal: find a split (k, τ) that minimizes an impurity measure at the leaves

  – Find best feature to split on
  – Find best threshold

18 Random Forests (RF)

– RF correct the overfitting bad "habit" of DTs.
– Training: Learn M trees on different subsets (random) of training data
– Prediction: Average of prediction of each tree
– Variance of the model averaging:

$$z_1 = f_1 \quad \Rightarrow \quad z_M = \frac{1}{M} \sum_{i=1}^{M} f_i$$

$$V(z_1) = \sigma^2 \quad \Rightarrow \quad V(z_M) = \frac{1}{M}\sigma^2 + \rho \frac{M-1}{M}\sigma^2$$

– Variance reduction ratio:

$$\frac{V(z_1)}{V(z_M)} = \frac{M}{1 + \rho(M-1)}$$

– 2 techniques for decorrelating trees:

  – **Bagging**: Randomize training data
  – Randomized feature selection